

Ensemble Methods: Boosting vs. Bagging

An Empirical Study

AI Academy Final Project Report, Summer 2026

Shahin Alakparov, Narmina Ibragimova, Hasan Mammadov

AI Academy, National AI Center

Spring 2026

{sahin.alekperov25, narmine.ibrahimova25, hasan.memmedov25}@aiacademy.az

Abstract—We present a systematic empirical comparison of two ensemble learning paradigms—*boosting* and *bagging*—implemented entirely from scratch using NumPy. Our custom implementations include a CART Decision Tree (with $O(N \log N)$ split search), AdaBoost (discrete SAMME with decision stumps), and a Random Forest (bootstrap aggregation with random feature sub-sampling and OOB scoring). We evaluate on three benchmark datasets spanning low-dimensional medical data, high-dimensional image classification, and large tabular census data. Seven controlled experiments investigate scaling, noise robustness, bias-variance decomposition, and cluster structure. We find that *AdaBoost minimises Bias² more aggressively (0.0231 vs. 0.0317 for RF) while Random Forest achieves far lower Variance (0.0076 vs. 0.0131 for AdaBoost)*. Under label corruption at $\eta=0.20$, Random Forest degrades only marginally while AdaBoost suffers a significant accuracy collapse, confirming that re-weighting amplifies label noise. On clean data both ensembles substantially outperform a single tree, with Random Forest achieving the highest AUC-ROC on Breast Cancer (where all figures are reported). Our implementations match scikit-learn baselines within 1% absolute accuracy, confirming correctness. Code is available at <https://github.com/shahin1717/ml-project> (tag `v1.0-final`).

I. INTRODUCTION

Ensemble methods are among the most successful machine learning techniques in practice [1]. By combining multiple base learners, ensembles reduce generalisation error more effectively than any individual model. Two dominant strategies exist: *boosting*, which builds an additive model sequentially to reduce bias, and *bagging*, which trains models independently in parallel to reduce variance.

The central question of this project is: “*Under what conditions does boosting outperform bagging, and vice versa, and why?*” To answer this rigorously we implement all algorithms from scratch, eliminating the risk of attributing differences to implementation details rather than algorithmic properties. We deliberately avoid ensemble classes from scikit-learn (except as a reference baseline) and use only NumPy for all numerical operations.

This setting is representative of real-world ML engineering: when a team needs to port an algorithm to a new language or hardware target, understanding the mathematical core—rather than relying on library abstractions—is essential.

Contributions. (1) Fully from-scratch implementations of CART, AdaBoost, and Random Forest with 93% test coverage. (2) Seven reproducible experiments with fixed random seeds on three diverse datasets. (3) Quantitative bias-variance decomposition comparing boosting and bagging. (4) Bonus implementations of Gradient Boosting and SAMME.R with comparative analysis.

The paper is organised as follows. Section II reviews the theoretical background. Section III describes our implementations. Section IV details the experimental design. Section V presents empirical results with analysis. Section VI synthesises findings, and Section VII concludes.

II. RELATED WORK

A. Decision Trees

CART [2] builds binary trees by greedily selecting splits that minimise impurity. The two standard impurity measures are Gini impurity $I_G = 1 - \sum_c p_c^2$ and Shannon entropy $I_E = -\sum_c p_c \log_2 p_c$. Trees are prone to overfitting on noisy data and have high variance—the primary motivation for ensemble methods.

B. AdaBoost

Freund and Schapire [3] showed that any *weak learner* (accuracy slightly above 50%) can be boosted into a strong learner. AdaBoost maintains a distribution over training samples, upweighting misclassified examples at each round. The final classifier is a weighted majority vote. Hastie *et al.* [1] proved that AdaBoost minimises an exponential loss and can be interpreted as gradient boosting in function space.

C. Bagging and Random Forests

Breiman [4] introduced bootstrap aggregation (bagging) as a general variance-reduction technique. Random Forests [5] extend bagging by sub-sampling features at each node split, which decorrelates individual trees and dramatically reduces variance beyond what pure bagging achieves. Out-of-bag (OOB) error provides an unbiased generalisation estimate without requiring a separate validation set.

D. Bias-Variance Trade-off

The expected prediction error decomposes as $\text{Err} = \text{Bias}^2 + \text{Variance} + \sigma_\epsilon^2$. Boosting primarily reduces bias by correcting errors iteratively; bagging primarily reduces variance by averaging [1]. Breiman [4] proved that averaging reduces variance by a factor proportional to the correlation between trees.

E. Unsupervised Methods

PCA [6] reduces dimensionality by projecting onto directions of maximum variance. K-Means [7] partitions data by minimising within-cluster sum of squares. DBSCAN [8] identifies clusters of arbitrary shape without requiring the number of clusters as input, labelling low-density points as noise.

III. METHODS

A. Decision Tree (CART)

We implement binary splitting with exhaustive threshold search over continuous features. For each feature we sort the N samples ($O(N \log N)$) and sweep all midpoints between consecutive distinct values using `np.cumsum` to compute weighted class counts incrementally-avoiding an inner Python loop over samples. The per-feature complexity is thus $O(N \log N + N)$ rather than $O(N^2)$. The overall split search is $O(PN \log N)$ where P is the number of candidate features.

The impurity reduction (information gain) for a candidate split is:

$$\Delta I = I(S) - \frac{|S_L|}{|S|} I(S_L) - \frac{|S_R|}{|S|} I(S_R) \quad (1)$$

Stopping criteria: (i) maximum depth reached, (ii) node samples $< \text{min_samples_split}$, (iii) pure node, (iv) all feature vectors identical. Feature importances are computed as normalised mean decrease in impurity (MDI) weighted by node sample count.

B. AdaBoost (Discrete SAMME)

We implement the discrete SAMME variant [1], which generalises AdaBoost to K -class problems. Let $w_i^{(1)} = 1/N$. At each round m :

- 1) Fit a stump h_m on weighted data.
- 2) Compute weighted error $\epsilon_m = \sum_i w_i^{(m)} \mathbf{1}[h_m(x_i) \neq y_i]$.
- 3) Compute estimator weight: $\alpha_m = \beta \left(\ln \frac{1-\epsilon_m}{\epsilon_m} + \ln(K-1) \right)$ where β is the learning rate.
- 4) Update and renormalise weights: $w_i^{(m+1)} \propto w_i^{(m)} \exp(\alpha_m \mathbf{1}[h_m(x_i) \neq y_i])$.

Early stopping triggers when $\epsilon_m \geq (K-1)/K$ (worse than random guessing). Each stump receives a seed derived from the master random state to ensure reproducibility.

C. Random Forest

Bootstrap sampling draws N samples with replacement per tree; the out-of-bag set is the complement. At each node, $\lfloor \sqrt{P} \rfloor$ features are randomly sampled before finding the best split. Trees are trained in parallel using `multiprocessing.Pool`. OOB predictions are accumulated only from trees for which a sample was *not* used during training, giving an unbiased score $\hat{\text{acc}}_{\text{OOB}}$. Feature importances average the MDI scores across all trees.

D. Unsupervised Pipeline

PCA centres the data and decomposes the covariance matrix via `np.linalg.eigh`. **K-Means** uses Lloyd's algorithm with k -distance initialisation over multiple restarts, tracking the lowest inertia. **DBSCAN** identifies core points (at least `min_samples` neighbours within ϵ) and expands clusters via BFS.

IV. EXPERIMENTAL SETUP

A. Datasets

Table I summarises the three datasets. Each presents different modelling challenges.

TABLE I
DATASET CHARACTERISTICS AFTER PREPROCESSING.

Dataset	N	P	Classes	Challenge
Breast Cancer	569	30	2	Small, clean
Adult Income	45 222	100	2	Large, imbalanced
MNIST 3-vs-8	6 000	784	2	High-dim, visual

Breast Cancer Wisconsin is a small, clean binary dataset predicting malignancy from 30 real-valued cell-nucleus measurements. The low noise and clean boundary make it an ideal benchmark.

Adult Income predicts whether annual income exceeds \$50K. The dataset is class-imbalanced ($\approx 76\%$ negative class), contains mixed continuous and categorical features (one-hot encoded to 100 columns after one-hot expansion of 8 categorical variables). Rows with missing values (marked ?) were discarded via `dropna`, reducing the dataset size from 48,842 to 45,222 cleaned examples.

MNIST 3-vs-8 is a 784-dimensional binary classification problem. Digits 3 and 8 share curved strokes, making them one of the hardest MNIST binary pairs. We subsample 6,000 images (≈ 3000 per class) to satisfy the specification requirement of ≥ 5000 samples while keeping computation manageable.

B. Preprocessing

Features were standardised (zero mean, unit variance) using training statistics only, applied to the test set identically. No data augmentation was applied.

C. Evaluation Protocol

Experiment 1 (correctness) uses a single fixed 80/20 train-test split. Experiments 2–3 (scaling curves) use a single fixed split to enable staged ensemble predictions across estimator counts. Experiment 4 (head-to-head) uses stratified 5-fold cross-validation with `random_state=42`. Experiment 5 (noise robustness) uses a single fixed split, varying only the label corruption fraction. We report accuracy, macro F1-score, and AUC-ROC. Experiment 6 (bias-variance) uses $B=100$ bootstrap replicates of the Breast Cancer training set, evaluating all models on a fixed held-out test split. All results reported below used 100 estimators for both ensembles.

V. RESULTS

A. Experiment 1: Correctness Verification

We compare our custom `DecisionTree` against `sklearn.tree.DecisionTreeClassifier` (Table II). Crucially, all differences are within 1%, satisfying the $\leq 2\%$ threshold required by the specification.

TABLE II
EXPERIMENT 1: BASELINE ACCURACY-CUSTOM VS. SKLEARN.

Dataset	Custom	Sklearn	Stump	$ \Delta $
Breast Cancer	0.9469	0.9469	0.8850	0.0000
Adult Income	0.8092	0.8105	0.7543	0.0013
MNIST 3-vs-8	0.9292	0.9317	0.8392	0.0025

Why is Adult accuracy “only” 81%? Adult Income is genuinely difficult: the negative class accounts for 76% of samples, the feature space is high-dimensional after OHE (100 columns), and the income boundary is a complex non-linear function of age, education, occupation, and capital gains. A single unpruned tree achieves 81% partly by defaulting toward the majority class. The custom tree matches sklearn within 0.2% (0.8092 vs. 0.8105), well within the $\leq 2\%$ threshold. The stump at 75% barely exceeds the naïve majority-class baseline ($\approx 76\%$).

Why does the stump achieve only $F1=0.43$ on Adult? A depth-1 tree can only partition the 100-dimensional space with one binary cut. On a balanced dataset this is already weak; on an imbalanced one it tends to predict the majority class for most observations, which yields high accuracy but very low recall for the minority, collapsing macro F1.

MNIST accuracy of 93% is impressive for a single unpruned tree on raw pixels. Digits 3 and 8 share curved strokes and are one of the hardest MNIST binary pairs, yet 784 pixel features provide enough discriminative signal for a deep tree. Our custom tree achieves 92.9% vs. sklearn’s 93.2%, a difference of 0.25%.

B. Experiment 2: AdaBoost Scaling

Figure 1 shows test and train accuracy as a function of the number of stumps T .

Observation. On Breast Cancer, test accuracy plateaus near 0.97 by $T \approx 30$ stumps; additional estimators do not help. On




Fig. 1. AdaBoost train/test accuracy vs. number of stumps on Breast Cancer. Train accuracy quickly reaches 1.0; test accuracy plateaus around $T=30$.

Adult Income (Fig. 2), convergence is slower because the boundary is more complex, and the training accuracy reaches 1.0 at $T \approx 80$.




Fig. 2. AdaBoost scaling on Adult Income. Slower convergence due to high-dimensional, imbalanced feature space.

Why does AdaBoost not overfit here? For exponential loss and stumps, Schapire showed that the margin distribution keeps improving even after training error reaches zero, explaining the continued test accuracy improvement before the plateau.

C. Experiment 3: Random Forest Scaling

Number of trees (Fig. 3). Accuracy is noisy for $T < 20$ trees (high variance from limited averaging) then smooths out above $T=50$. Beyond $T=100$ there is negligible gain-variance reduction saturates as the law of large numbers kicks in.

Tree depth (Fig. 4). Shallow trees (depth ≤ 4) underfit; deep trees (depth ≥ 15) memorise training data without proportional test improvement. The sweet spot is depth ≈ 10 for the Breast Cancer dataset. This contrasts with AdaBoost: stumps (depth 1) work for AdaBoost because the ensemble aggregates thousands of complementary cuts.




Fig. 3. Random Forest accuracy vs. number of trees (Breast Cancer). Test accuracy stabilises near $T=50$.




Fig. 4. Random Forest accuracy vs. max tree depth (Breast Cancer). Accuracy peaks around $\text{max_depth}=10$; deeper trees overfit.

D. Experiment 4: Head-to-Head Comparison

Figures 5–7 show box plots from 5-fold cross-validation at $T=100$ estimators.

Key findings. (1) Both ensembles reduce inter-quartile range significantly compared with the single tree, confirming variance reduction. (2) On Breast Cancer, AdaBoost achieves the highest mean accuracy (0.9755 ± 0.0186) while Random Forest is close behind (0.9631 ± 0.0203); on AUC-ROC both reach > 0.98 . (3) On Adult Income, our custom Random Forest (0.8623 ± 0.0032) marginally outperforms sklearn’s RF (0.8507 ± 0.0023), demonstrating that the custom implementation is competitive. (4) On MNIST with 784 raw-pixel features, Random Forest dominates with 0.9808 ± 0.0035 accuracy vs. AdaBoost’s 0.9505 ± 0.0032 —the random feature subsampling effectively focuses each tree on the most informative pixel regions.

E. Experiment 5: Noise Robustness

We randomly flip a fraction η of training labels.

Why does AdaBoost collapse under noise? The re-weighting mechanism assigns exponentially higher weight to




Fig. 5. 5-fold CV accuracy on Breast Cancer. Both ensembles substantially outperform the single tree with lower variance.




Fig. 6. 5-fold CV AUC-ROC on Breast Cancer. Random Forest achieves the highest median AUC-ROC with the smallest inter-quartile range.

misclassified points. When some of those points are mislabelled (genuinely hard boundary samples that are now labelled incorrectly), the stump trained in the next round attempts to classify what are in fact random-label points, degrading generalisation. At $\eta=0.20$, approximately one fifth of all re-weighted points carry corrupted labels, causing the ensemble to memorise noise.

Why does Random Forest stay stable? Each tree in the forest is trained on a separate bootstrap sample, so each individual tree is only partially exposed to the corrupted labels. When predictions are averaged, the impact of any single corrupted tree is diluted. Furthermore, since Random Forest does not focus more resources on hard samples the way AdaBoost does, it avoids the amplification effect. This amplification is most pronounced on Breast Cancer, where the class boundary is clean; on Adult Income, where the boundary is already noisy and complex, AdaBoost’s re-weighting degrades more gracefully than on cleaner data, yielding a smaller accuracy




Fig. 7. 5-fold CV accuracy on MNIST 3-vs-8. High dimensionality benefits both ensemble methods equally.




Fig. 8. Accuracy vs. label noise fraction η on Breast Cancer. Random Forest degrades gracefully; AdaBoost collapses at high noise.

drop than Random Forest at $\eta=0.20$.

F. Experiment 6: Bias-Variance Decomposition

Using $B=100$ bootstrap training sets, we estimate Bias^2 and Variance following Breiman [4] (Table III and Fig. 10).

TABLE III
BIAS-VARIANCE DECOMPOSITION (BREAST CANCER, $B=100$).

Model	Bias^2	Variance
Single Tree	0.0408	0.0352
AdaBoost	0.0231	0.0131
Random Forest	0.0317	0.0076

Interpretation. The single tree has the highest values on both dimensions: it overfits some bootstrap samples (high variance) while still making systematic boundary errors (high bias).

AdaBoost reduces Bias^2 by 43% relative to the tree (from 0.0408 to 0.0231) by iteratively correcting systematic errors.




Fig. 9. Noise robustness on Adult Income. Random Forest retains a higher absolute accuracy at every η , but AdaBoost degrades *less* in relative terms ($\Delta=0.0019$ vs. RF's $\Delta=0.0056$ from $\eta=0$ to 0.20)-the opposite of the degradation-rate ordering seen on Breast Cancer.




Fig. 10. Bias^2 and Variance for three models on Breast Cancer. AdaBoost wins on bias; Random Forest wins on variance.

Its Variance is moderate (0.0131) because each round's stump is deterministic given the weights, but the overall ensemble is sensitive to which samples are upweighted, introducing some cross-replicate variation.

Random Forest achieves the lowest Variance (0.0076, a 78% reduction from the single tree) through averaging of decorrelated trees. Its Bias^2 (0.0317) is intermediate: each tree is unbiased (full-depth), but feature sub-sampling means individual splits may not always choose the globally optimal feature, leaving some systematic error.

Total error. AdaBoost: $0.0231 + 0.0131 = 0.0362$. Random Forest: $0.0317 + 0.0076 = 0.0393$. On this dataset AdaBoost wins marginally on total decomposed error.

G. Experiment 7: Unsupervised Analysis

PCA. The first 7 of 30 principal components account for $\geq 90\%$ of the variance (Fig. 11). This moderate effective dimensionality explains why even a single unpruned tree achieves over 94% accuracy-the majority of 30 features are




Fig. 11. PCA scree plot on Breast Cancer. 90% of variance is captured by the first 7 principal components.




Fig. 12. K-Means elbow plot. Inertia decreases sharply up to $k=2$, confirming the binary class structure.

correlated cell-nucleus measurements, and the 7-PC projection captures the essential discriminative structure.

K-Means. The elbow at $k=2$ (Fig. 12) aligns with the binary classification task. The K-Means clusters recover the malignant/benign separation with an Adjusted Rand Index (ARI) of **0.6536**. The moderate ARI reflects that the two classes partially overlap in PCA space (Fig. 13), which is the same overlap that challenges classifiers.

DBSCAN. With $\varepsilon=3.5$ (chosen from the k-distance plot, Fig. 14), DBSCAN achieves ARI of 0.0653, substantially lower than K-Means. This low ARI indicates that density-based clustering struggles on the Breast Cancer dataset: the two classes are not separated by a density gap but rather blend along a continuous manifold. DBSCAN additionally labels approximately 13.2% of samples as noise-boundary-region points corresponding to the clinically ambiguous cases that are hardest to classify correctly.

H. Bonus: SAMME vs. SAMME.R

SAMME.R uses probability estimates from each stump rather than hard predictions, typically converging faster and yielding better-calibrated probabilities. Figure 15 confirms that




Fig. 13. 2-D PCA projection coloured by true labels (left), K-Means clusters (centre), and DBSCAN clusters (right).




Fig. 14. K-distance plot for DBSCAN parameter selection. The knee around $\varepsilon \approx 3.5$ guides our choice.

SAMME.R achieves lower log-loss at the same number of estimators.

I. Bonus: Gradient Boosting vs. AdaBoost

Our custom Gradient Boosting implementation fits trees on negative gradients of the log-loss. It outperforms AdaBoost on Breast Cancer (Fig. 16), converging to a higher accuracy plateau, because it uses deeper trees rather than stumps and directly optimises log-loss rather than exponential loss, which is more robust to near-boundary samples. On Adult Income the two methods are nearly tied (AdaBoost 0.8310 ± 0.0124 vs. GBM 0.8300 ± 0.0259), suggesting GBM’s advantage is most pronounced on cleaner, lower-dimensional data.

VI. DISCUSSION

A. When Does Boosting Win?

On *clean, small, and well-structured* datasets (Breast Cancer), AdaBoost matches or slightly exceeds Random Forest in total Bias²+Variance (0.0362 vs. 0.0393). Its sequential focus on hard samples drives the decision boundary into complex regions that naive averaging misses.

B. When Does Bagging Win?

Random Forest dominates in *three* scenarios:

- 1) **Label noise.** Averaging independent trees dilutes corrupted-label influence; AdaBoost amplifies it.
- 2) **High variance.** On smaller datasets (or noisier features) the bootstrapped averaging provides stability that sequential re-weighting cannot.




Fig. 15. Calibration comparison: SAMME (discrete) vs. SAMME.R (real) across datasets. SAMME.R converges faster and produces better-calibrated probabilities.




Fig. 16. GBM vs. AdaBoost test accuracy on Breast Cancer. Gradient Boosting converges to a higher plateau.

- 3) **High-dimensional feature spaces.** Random feature subsampling at each node effectively performs implicit feature selection per tree, making Random Forest naturally robust to irrelevant features without pre-processing.

C. Connection to Unsupervised Structure

The PCA analysis shows that Breast Cancer’s variance is spread across 7 principal components (not a tight 2D cluster), yet both ensembles achieve near-perfect AUC-ROC (≥ 0.98). The high K-Means ARI (0.6536) confirms that the malignant/benign boundary is largely linearly separable in PC space. By contrast, DBSCAN’s low ARI (0.0653) and noise fraction of 13.2% reveal that the density structure is diffuse—there is no clean density gap between classes. The 13.2% noise points are the boundary-region samples that AdaBoost re-weights heavily in later rounds, which explains both its lower bias and its vulnerability to label noise: it cares disproportionately about the hardest cases.

D. Implementation Correctness

Our custom tree matches sklearn within 1% absolute accuracy on all three datasets (Table II), confirming that the $O(N \log N)$ cumsum sweep correctly implements the CART criterion without floating-point errors.

VII. CONCLUSION

We implemented CART, AdaBoost (SAMME), and Random Forest from scratch and evaluated them across seven controlled experiments. The key finding is that neither algorithm uniformly dominates: *AdaBoost reduces Bias² more aggressively*; *Random Forest reduces Variance more dramatically and is significantly more robust to label corruption*. Practitioners should prefer boosting when data is clean and structured, and bagging when noise, class imbalance, or high dimensionality is a concern.

Limitations. Experiments use fixed hyperparameters (100 estimators, `max_features=sqrt`). A full grid search might narrow the performance gap. All experiments use the full datasets: Breast Cancer (569 samples), Adult Income (full 45,222 examples after dropping rows with missing values via `dropna`), and MNIST 3-vs-8 (6,000 samples, satisfying the ≥ 5000 specification).

REFERENCES

- [1] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning*, 2nd ed. Springer, 2009.
- [2] J. R. Quinlan, “Induction of decision trees,” *Machine Learning*, vol. 1, no. 1, pp. 81–106, 1986.
- [3] Y. Freund and R. E. Schapire, “A decision-theoretic generalization of on-line learning and an application to boosting,” *J. Comput. Syst. Sci.*, vol. 55, no. 1, pp. 119–139, 1997.
- [4] L. Breiman, “Bagging predictors,” *Machine Learning*, vol. 24, no. 2, pp. 123–140, 1996.
- [5] L. Breiman, “Random forests,” *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.
- [6] K. Pearson, “On lines and planes of closest fit to systems of points in space,” *Philosophical Magazine*, vol. 2, no. 11, pp. 559–572, 1901.
- [7] S. P. Lloyd, “Least squares quantization in PCM,” *IEEE Trans. Inf. Theory*, vol. 28, no. 2, pp. 129–137, 1982.
- [8] M. Ester, H.-P. Kriegel, J. Sander, and X. Xu, “A density-based algorithm for discovering clusters in large spatial databases with noise,” in *Proc. 2nd Int. Conf. Knowledge Discovery and Data Mining (KDD)*, 1996, pp. 226–231.